

6. Search Trees

Goals

- Understand search, insert, and delete in binary search trees
- Learn search, insert, and delete in red-black trees
- Learn search, insert, and delete in B-trees
- Understand multidimensional search trees

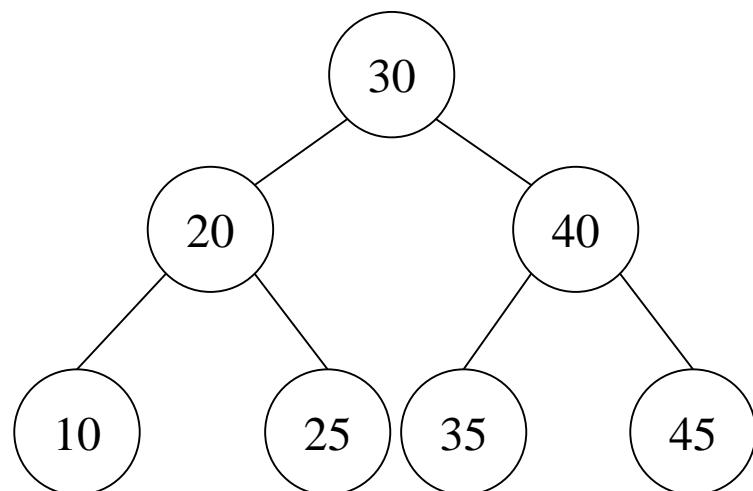
Record, Key, Search Tree

- Record
 - Storage unit including all information about an object
 - e.g., record for a person
 - Name, national ID, home address, mobile phone number, etc.
- Field
 - Each item in the record
 - e.g., name
- Key
 - Field(s) that represent records uniquely
 - Key may consist of one field or multiple fields.
- Search tree
 - Contains keys and record pointers
 - Provides searching mechanism by keys

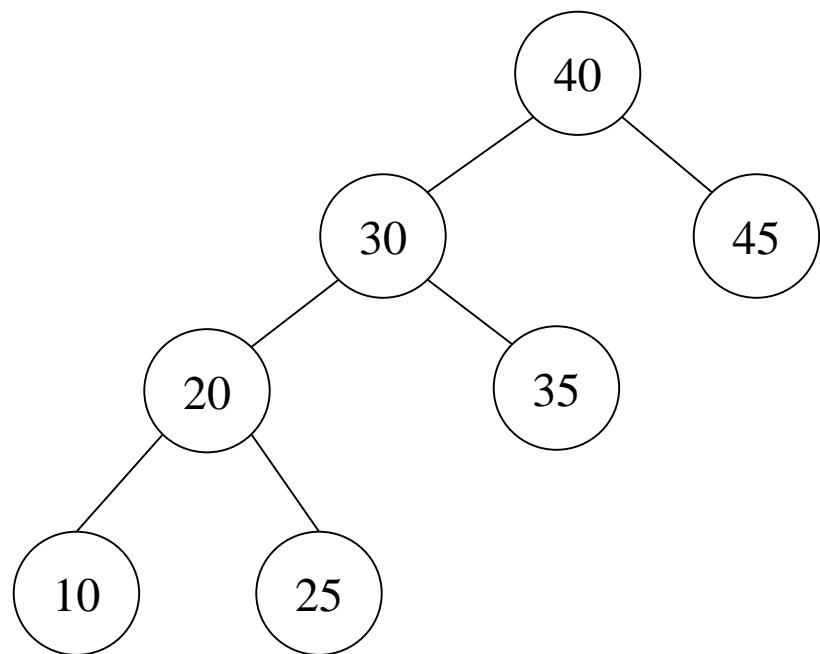
Binary Search Tree

- Each node of a binary search tree contains one key.
- Each node has at most two children. (binary tree)
- For any node x , key of $x >$ any key in x 's left subtree, and key of $x <$ any key in x 's right subtree. (search tree)

Binary Search Tree



(a)



(b)

Search in Binary Search Tree

treeSearch(t, x)

▷ t : tree node

▷ x : search key

{

if ($t=\text{NIL}$ **or** $\text{key}[t]=x$) **then return** t ;

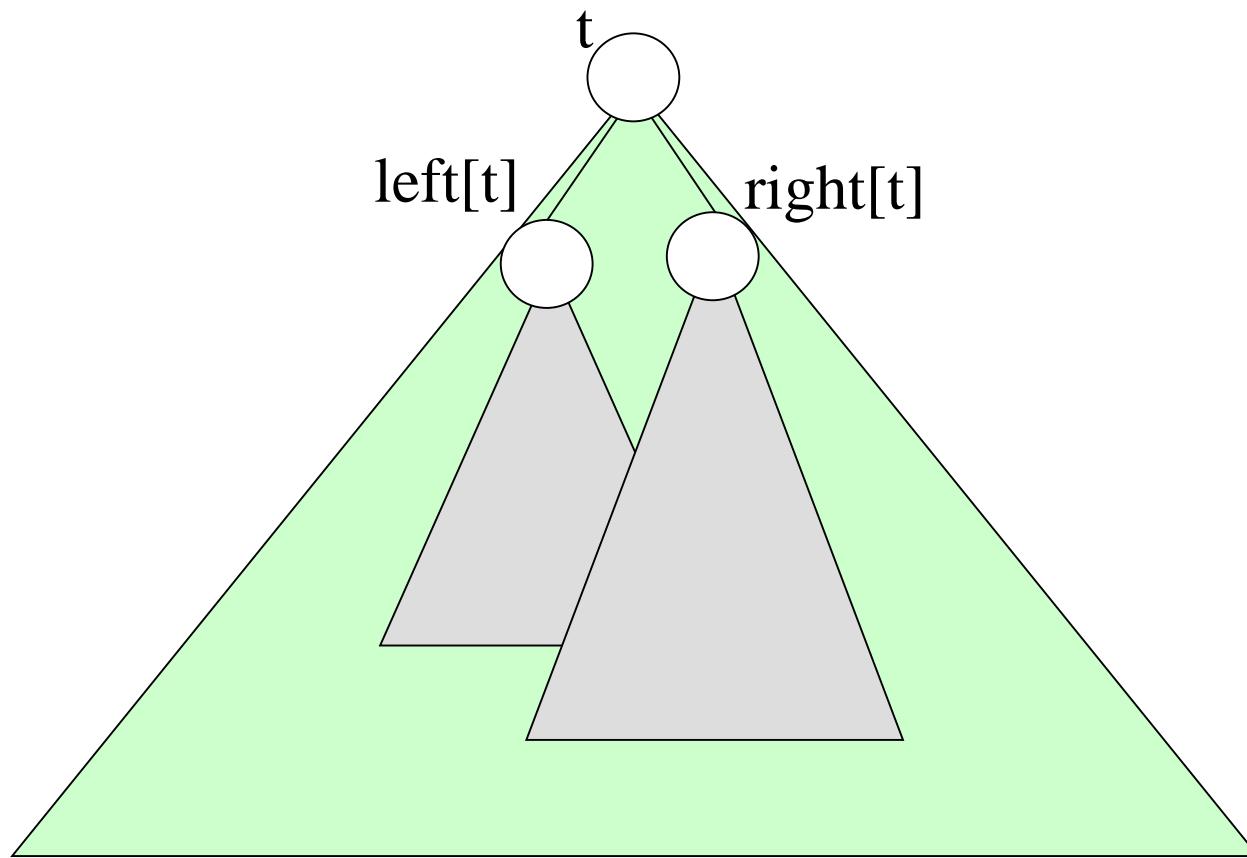
if ($x < \text{key}[t]$)

then return treeSearch(left[t], x);

else return treeSearch(right[t], x);

}

Recursion in Search



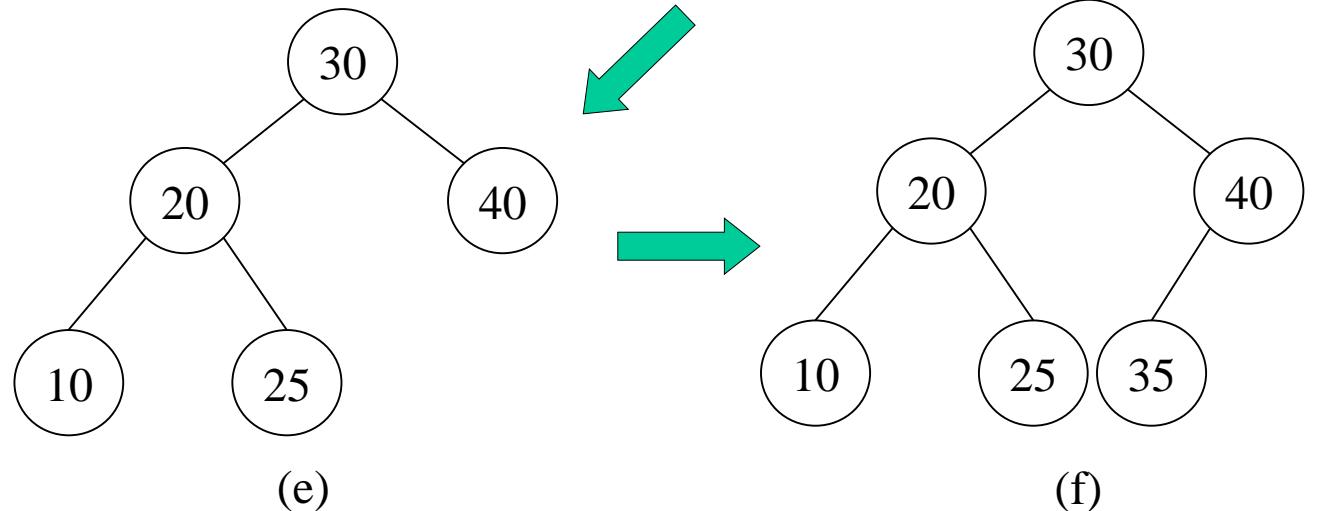
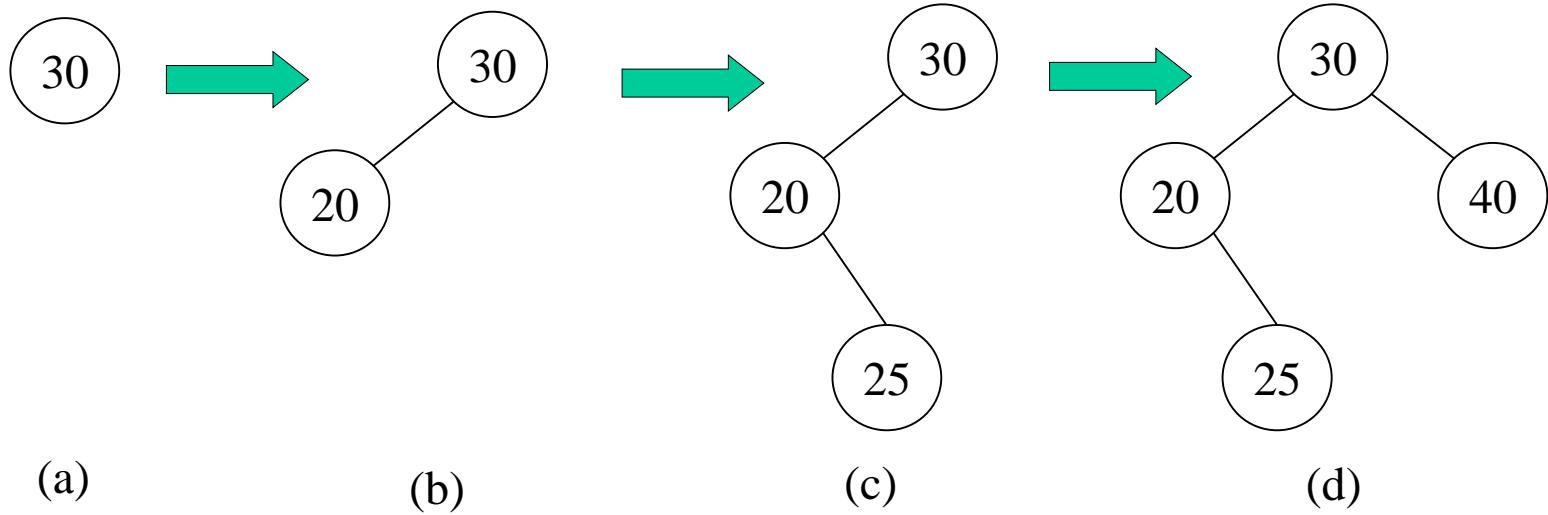
Insert in Binary Search Tree

treeInsert(t, x)

- ▷ t : tree node
- ▷ x : insert key
- ▷ return root after insertion

```
{  
    if (t=NIL) then {  
        key[r] ←  $x$ ; left[r] ← NIL; right[r] ← NIL;    ▷  $r$  : new node  
        return  $r$  ;  
    }  
    if ( $x < \text{key}(t)$ )  
        then {left[t] ← treeInsert(left[t],  $x$ ); return  $t$ ;}  
        else {right[t] ← treeInsert(right[t],  $x$ ); return  $t$ ;}  
}
```

Insert



Delete in Binary Search Tree

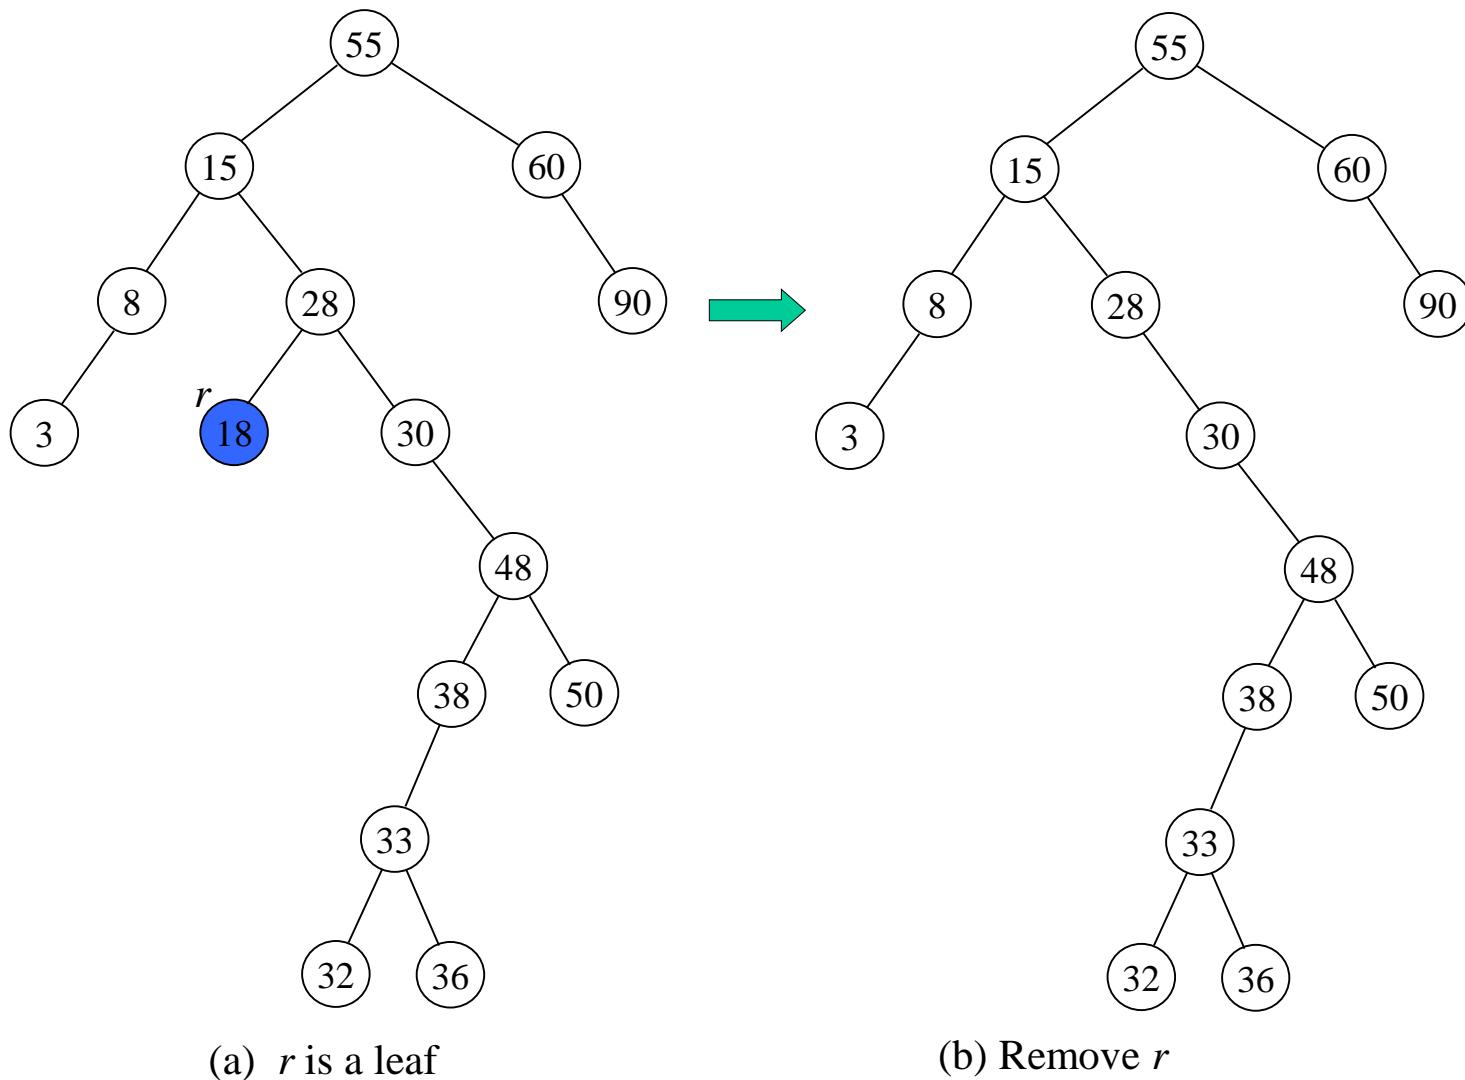
r : node to be deleted

- There are three cases
 - Case 1 : r is a leaf
 - Case 2 : r has one child
 - Case 3 : r has two children

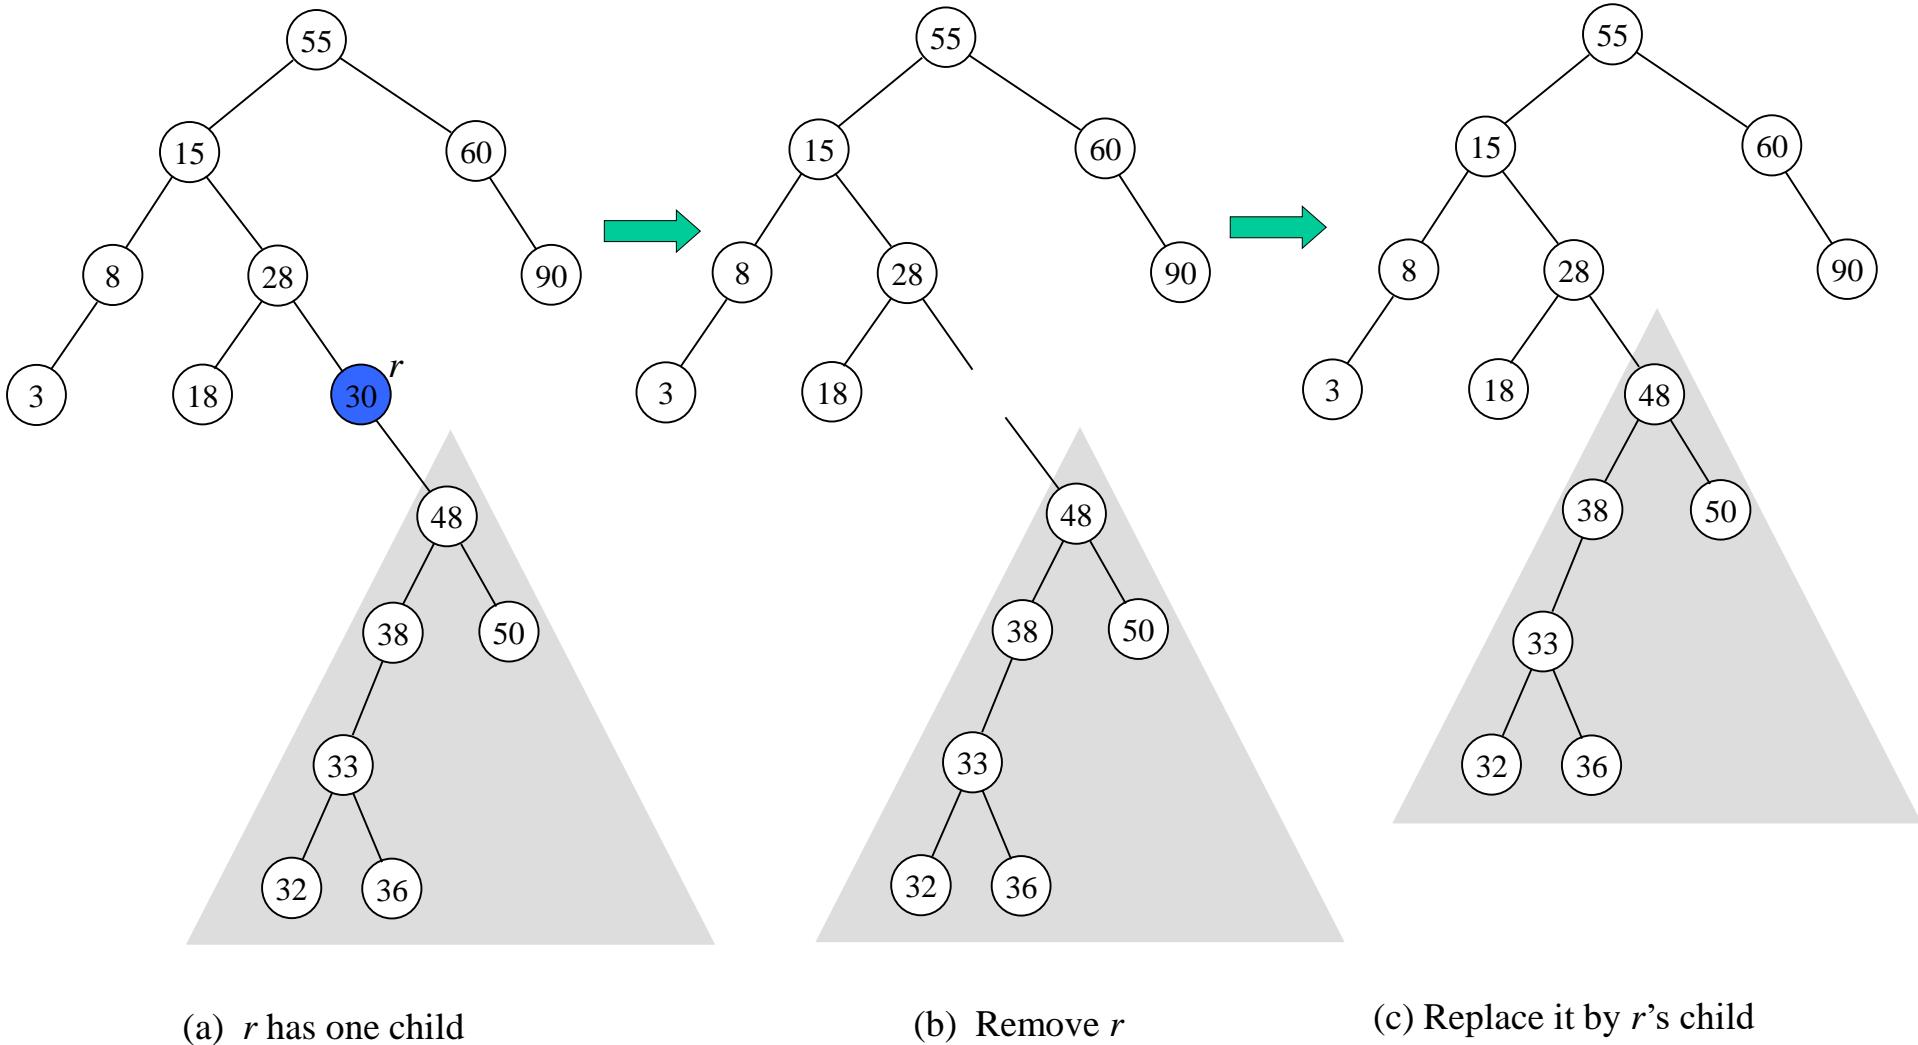
Delete

```
treeDelete( $t, r$ )
▷  $t$ : tree root
▷  $r$ : node to be deleted
{
    if ( $r$  is leaf) then                                ▷ Case 1
        remove  $r$ 
    else if ( $r$  has one child) then                  ▷ Case 2
        replace  $r$  by  $r$ 's child
    else                                            ▷ Case 3
        replace  $r$  by  $r$ 's successor  $s$  (smallest element in  $r$ 's right subtree)
        delete  $s$  (case 1 or 2)
}
```

Delete: Case 1



Delete: Case 2

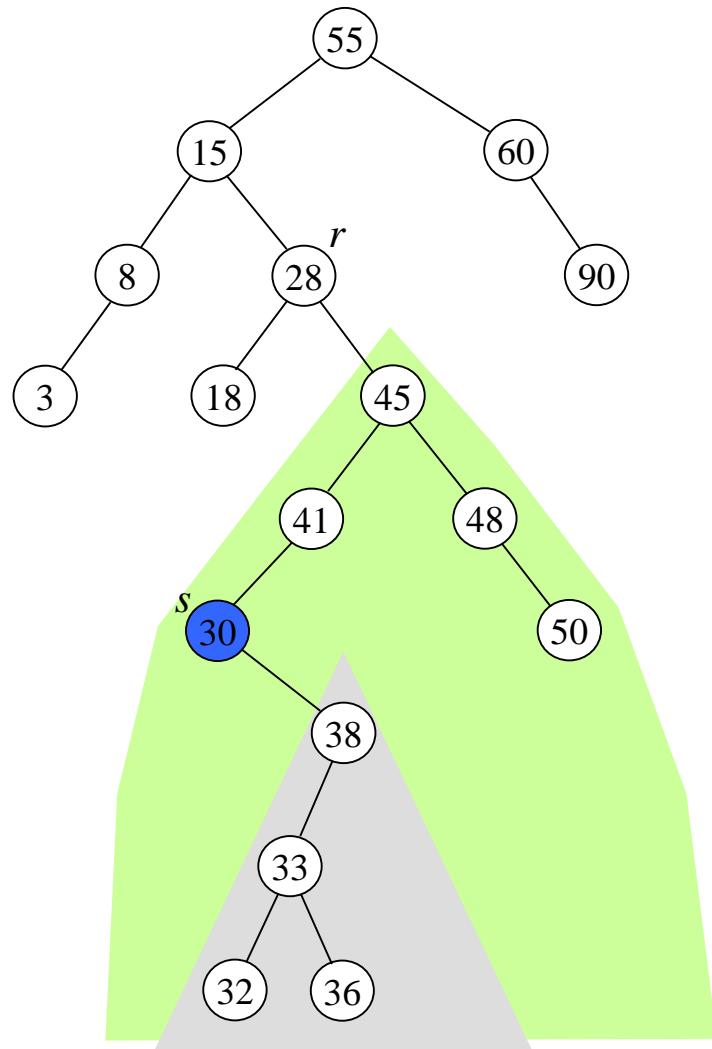


(a) r has one child

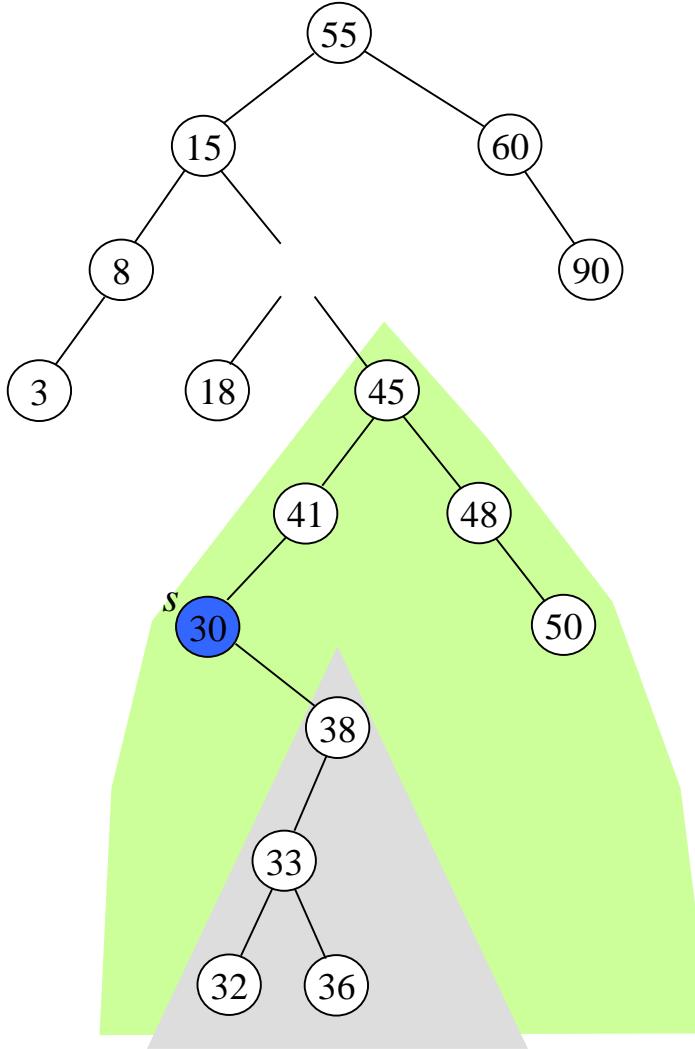
(b) Remove r

(c) Replace it by r 's child

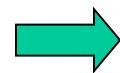
Delete: Case 3

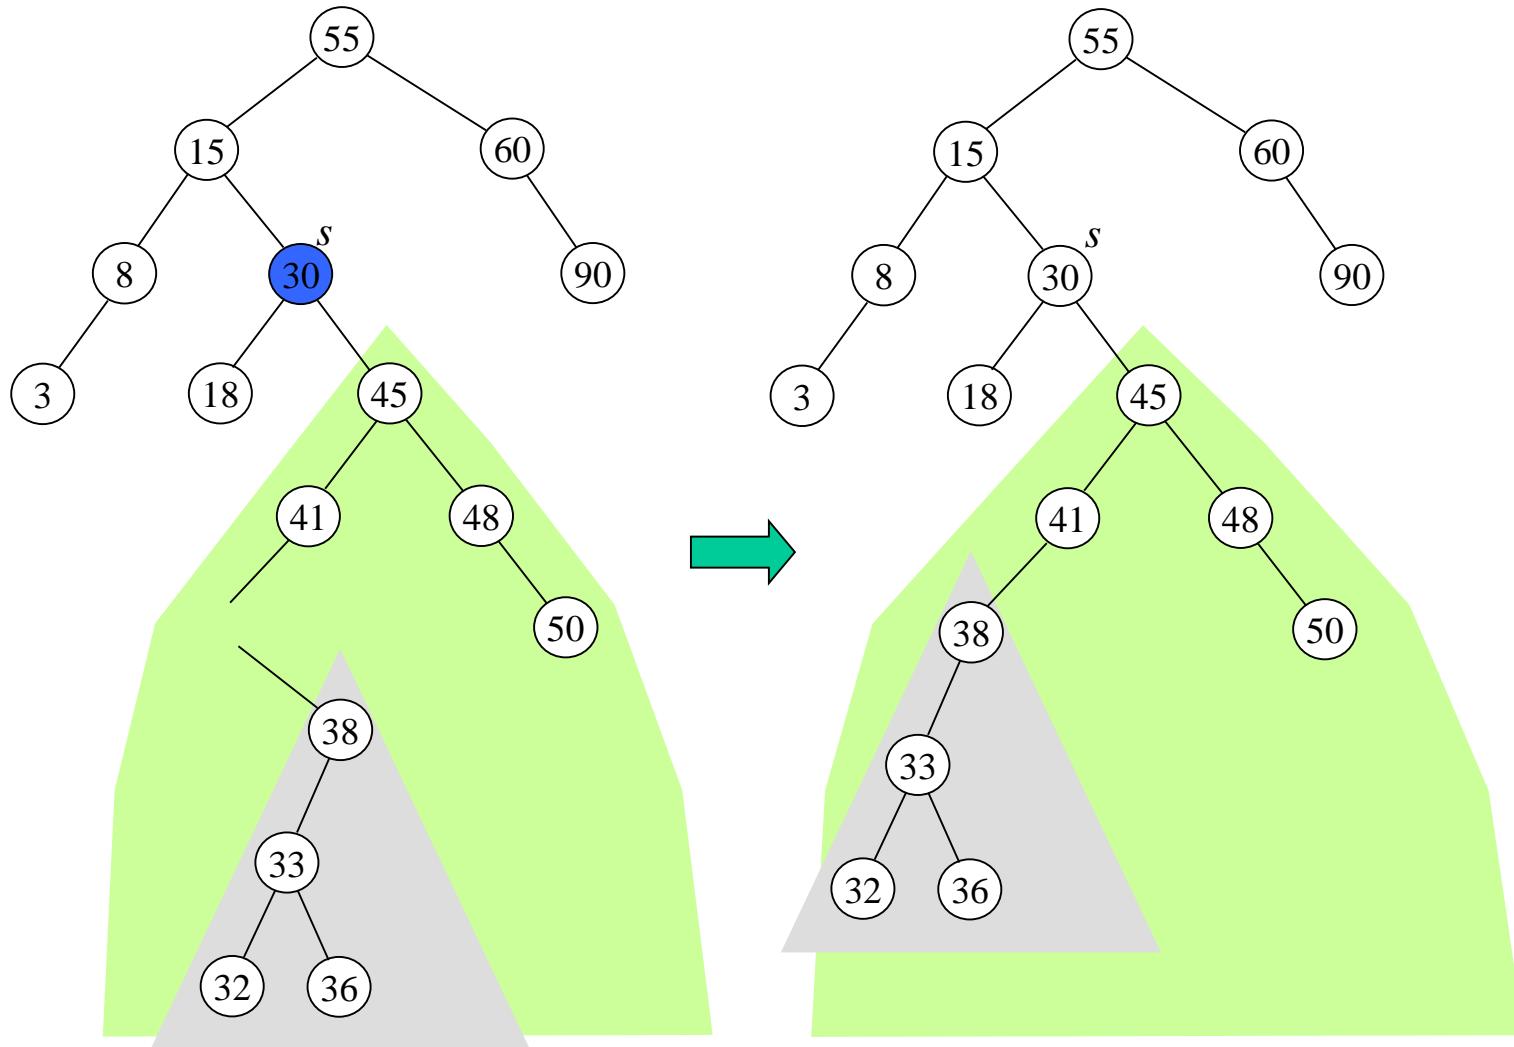


(a) Find r 's successor s



(b) Remove r

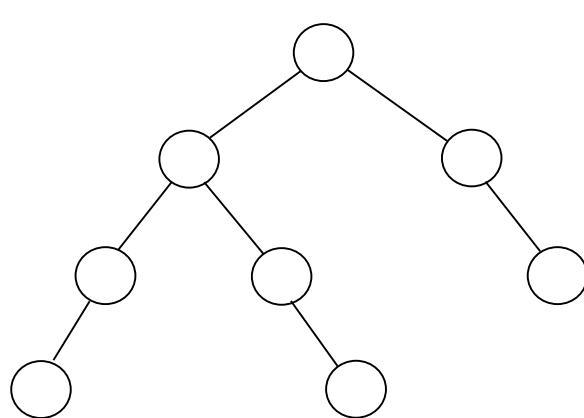




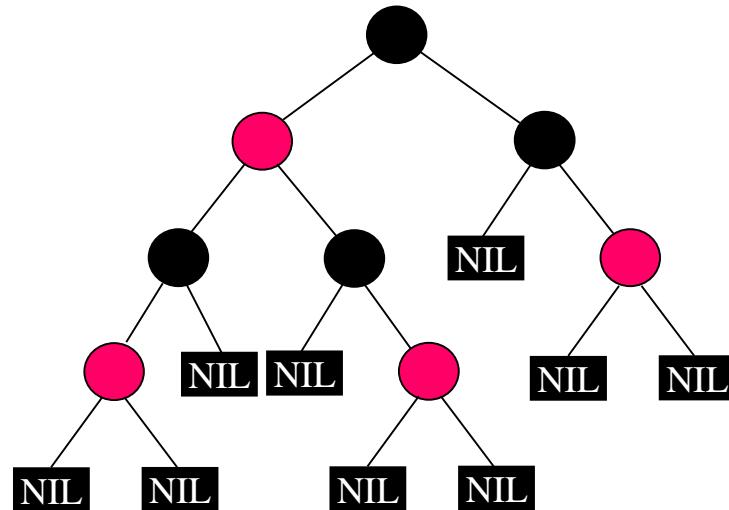
Red-Black Tree

- Red-black tree is a binary search tree with 1 bit information per node (color: red or black) satisfying the following red-black properties.
 - ① Root is black
 - ② Every leaf (NIL) is black
 - ③ If a node is red, both its children must be black
 - ④ All paths from root to leaf contain the same number of black nodes
- ✓ where each leaf has value NIL (actually each leaf has a pointer to a special node which has value NIL)

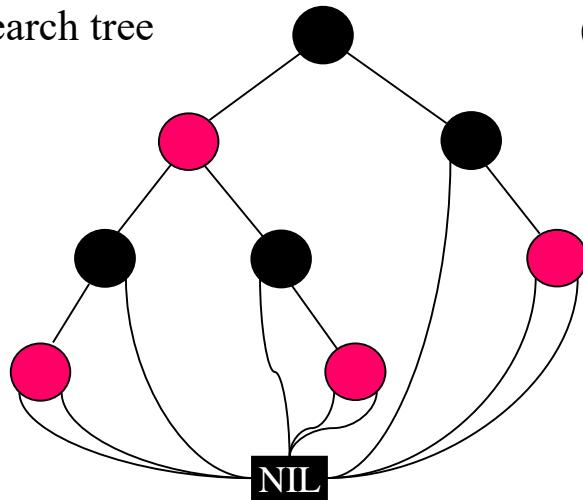
Red-Black Tree



(a) Binary search tree



(b) Red-black tree



(c) Actual implementation

Red-Black Tree

- Black height $bh(x)$: number of black nodes in the path from x to a leaf, excluding x
- A red-black tree has at least $2^{bh(r)} - 1$ internal nodes.
- A red-black tree with n internal nodes has height $h \leq 2 \log_2(n + 1)$.
 - $n \geq 2^{bh(r)} - 1$
 - $bh(r) \leq \log(n + 1)$
 - $h \leq 2bh(r) \leq 2 \log_2(n + 1)$

Search in Red-Black Tree

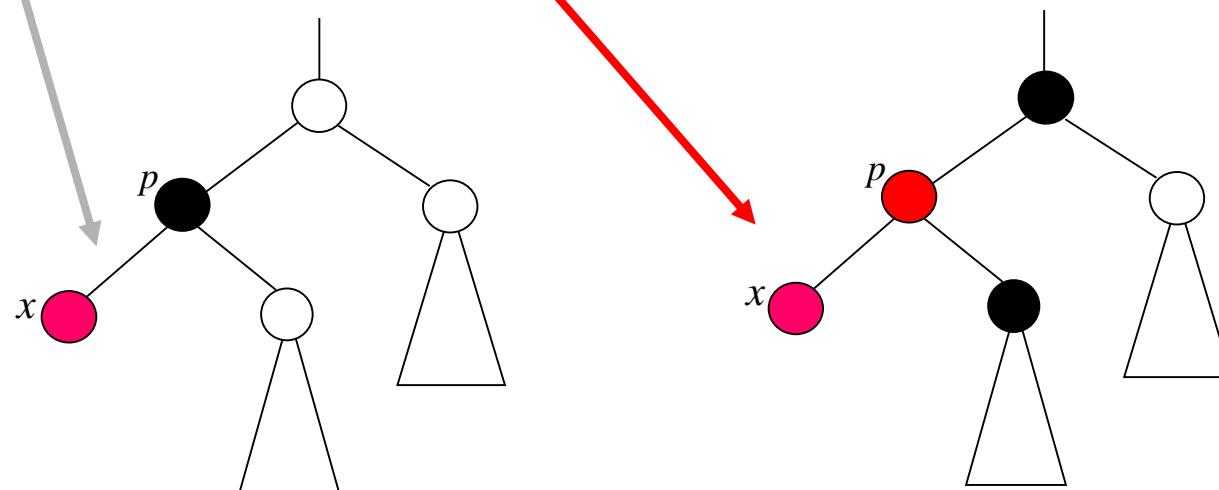
- does not change the tree
- procedure: same as that of binary search tree
- but in worst-case $O(\log n)$ time

Insert and Delete in Red-Black Tree

- modify the tree structure
- Basic operations: rotations
 - Left rotate
 - Right rotate
- need to change colors of some nodes to satisfy the red-black properties

Insert in Red-Black Tree

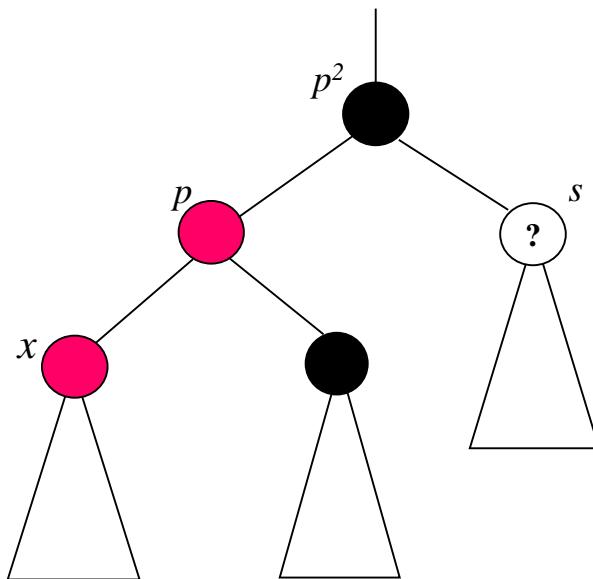
- Insert x : same as insert in binary search tree.
- Color x red.
- If the color of x 's parent p is
 - black, done.
 - red: property ③ is violated.



✓ Hence, consider the case p is red

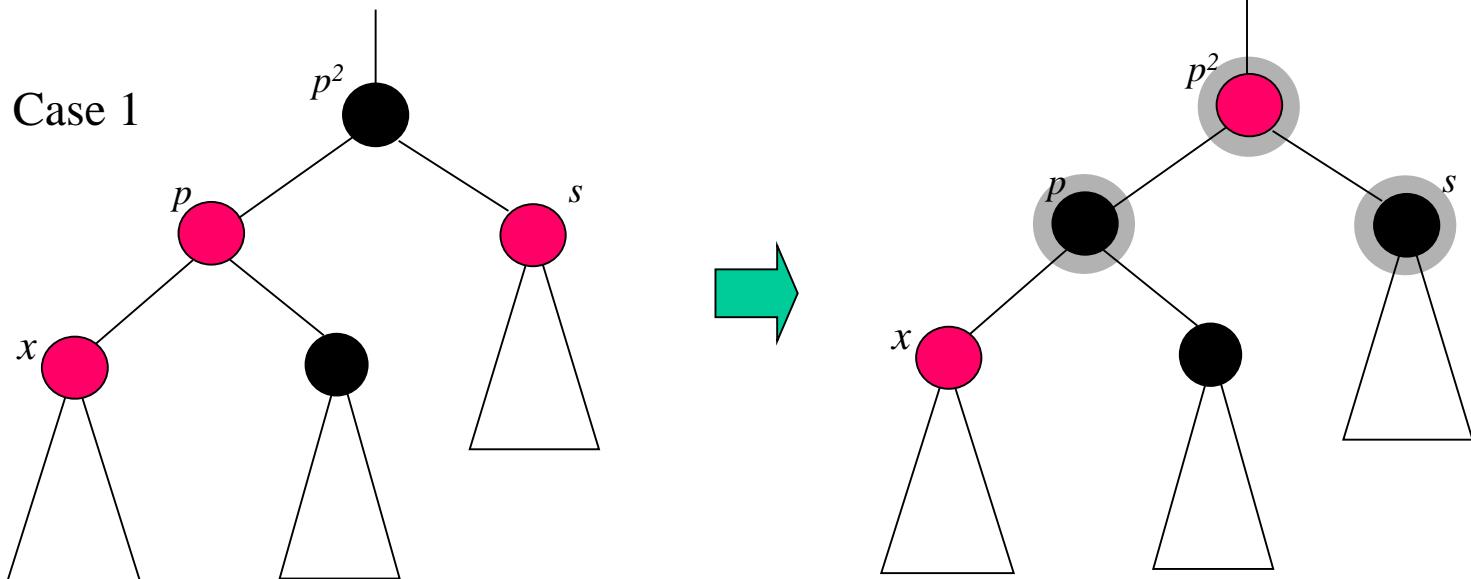
Insert in Red-Black Tree

- p^2 and x 's sibling must be black
- There are two cases depending on color of s (p 's sibling)
 - Case 1: s is red
 - Case 2: s is black



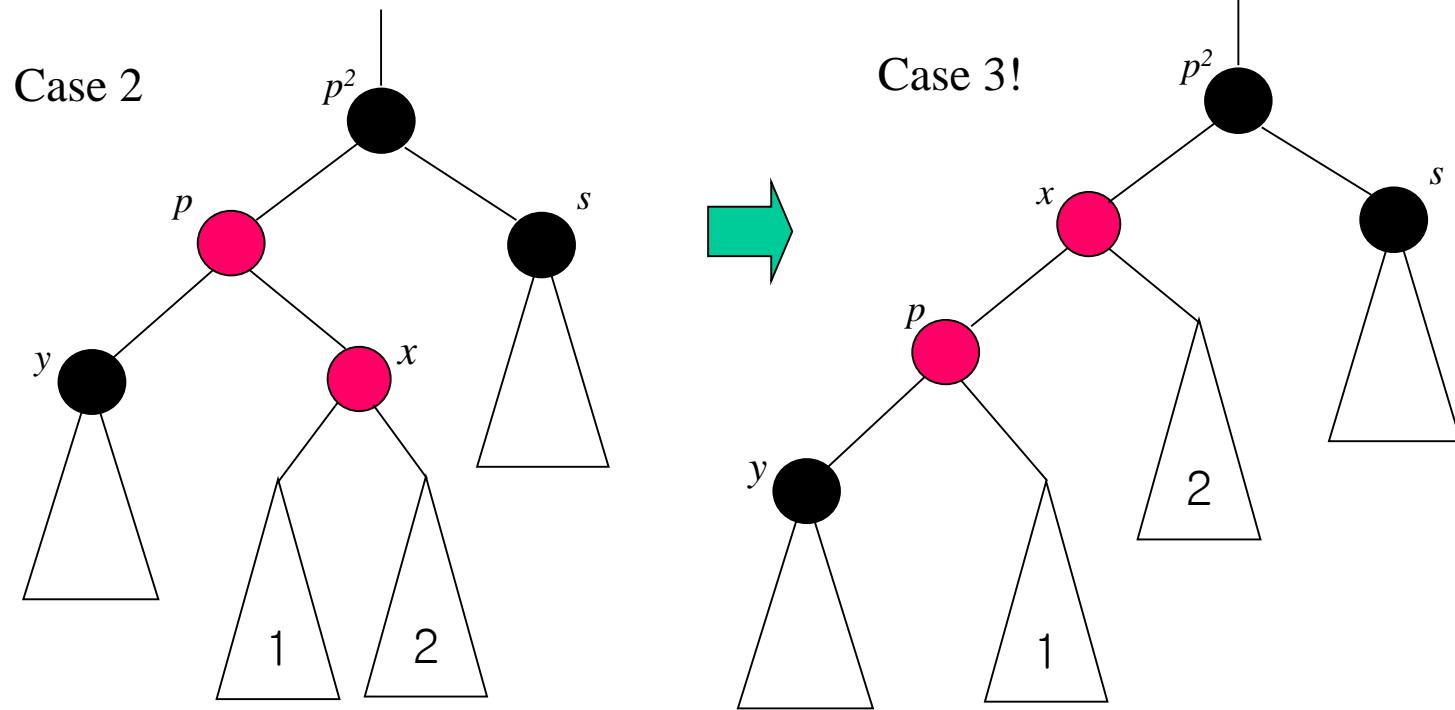
Case 1: s is red

● : color is changed



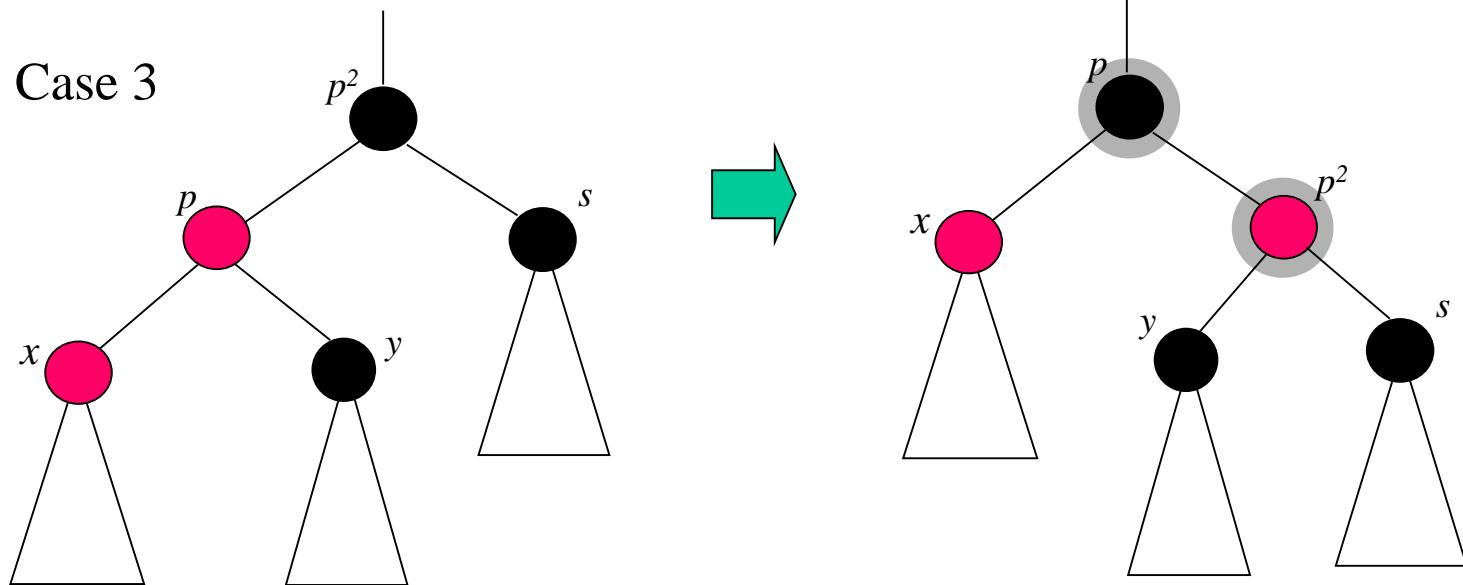
✓ p^2 may have the same problem

Case 2: s is black, x is p 's right child



Case 3: s is black, x is p 's left child

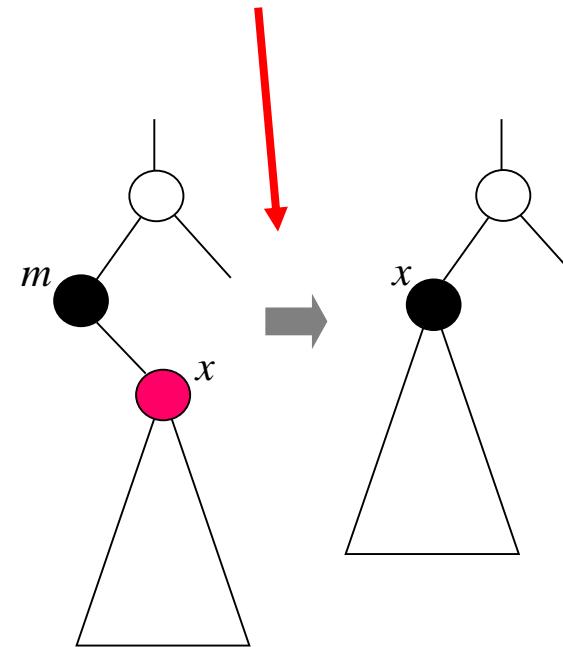
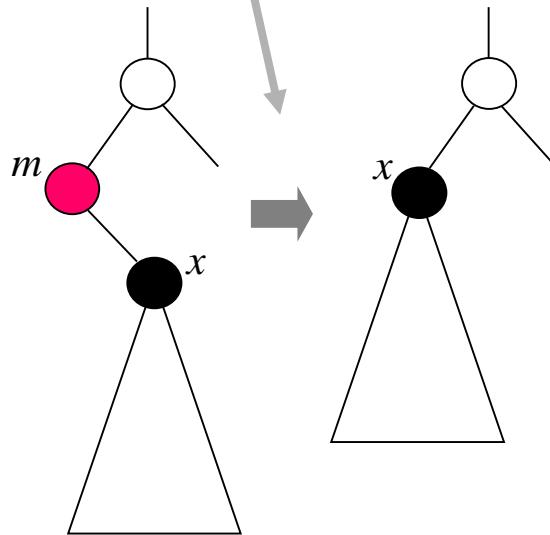
● : color is changed



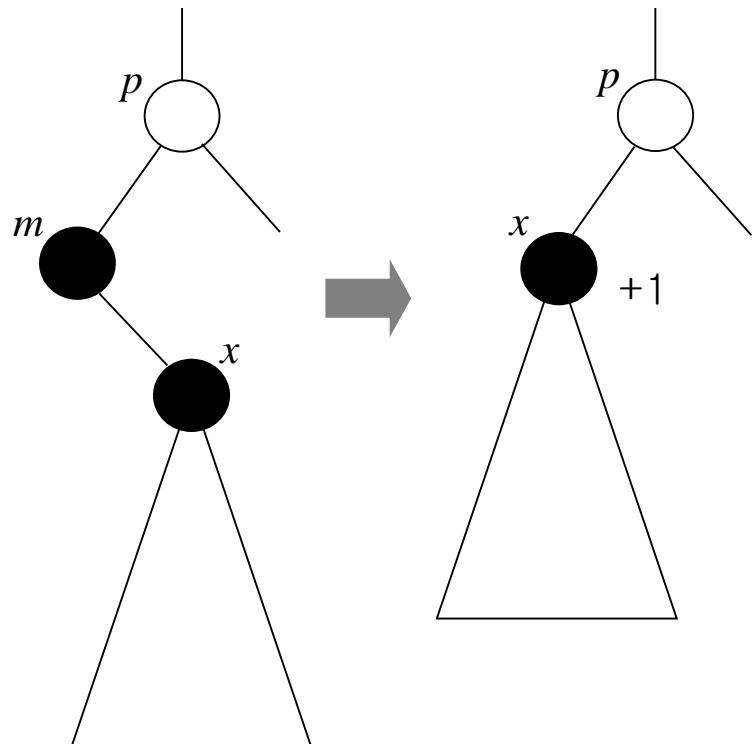
✓ insert done! $O(\log n)$ time

Delete in Red-Black Tree

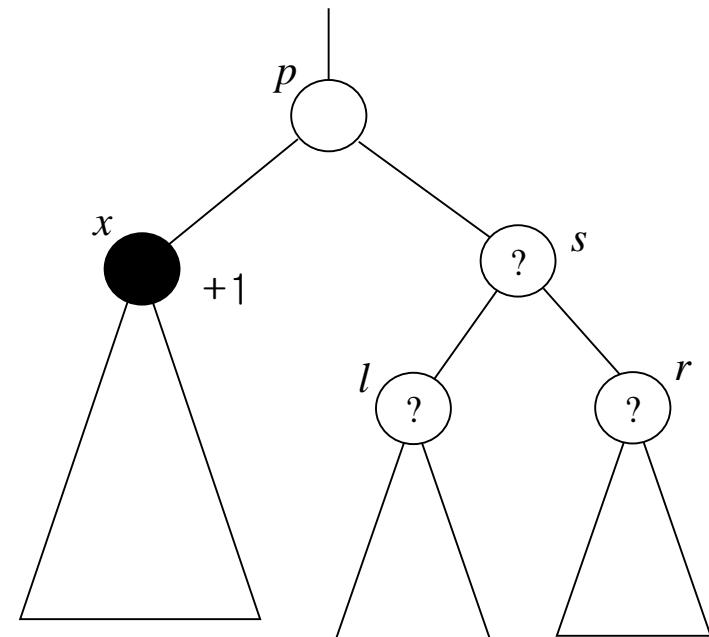
- Delete m : same as delete in binary search tree.
- Node m had at most one internal child.
- If m was red, we are done.
- If m was black and internal child was red, done.



✓ +1 means that x has an extra black

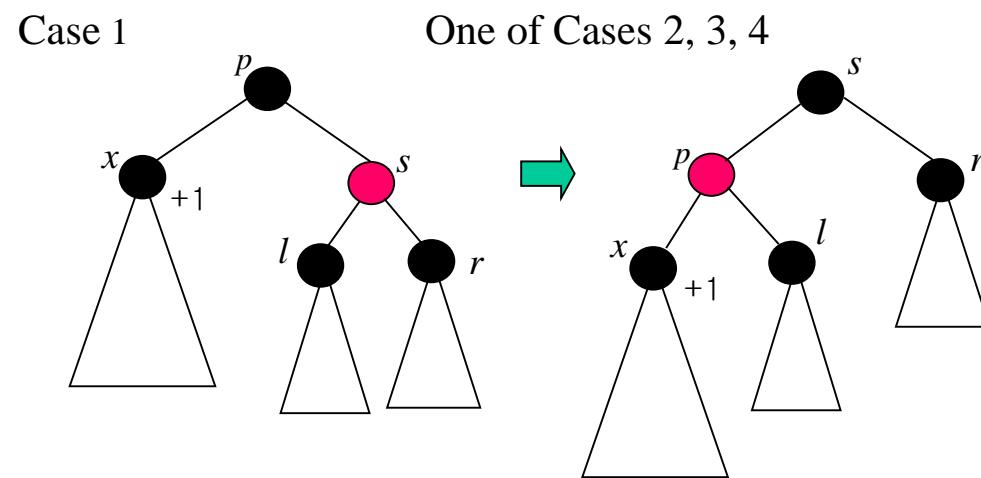


After m is deleted,
property ④ is violated

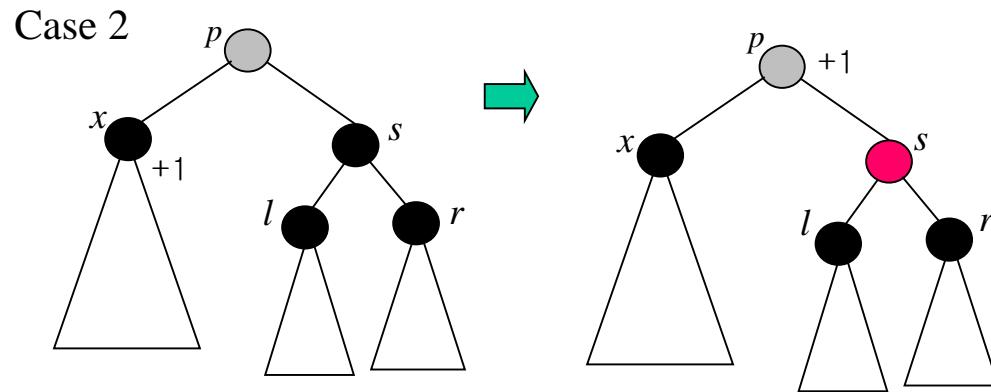


Cases by colors of s, l, r

Case 1: s is red

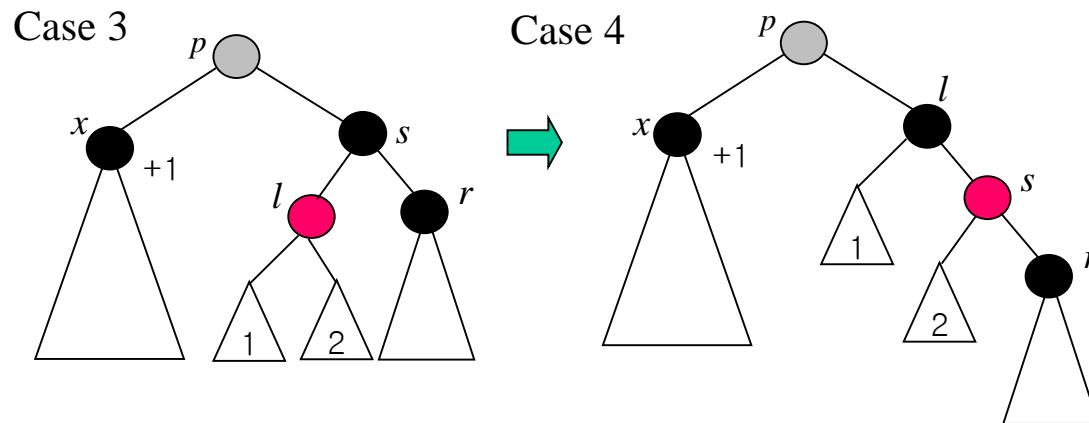


Case 2: s is black, and both of s 's children are black

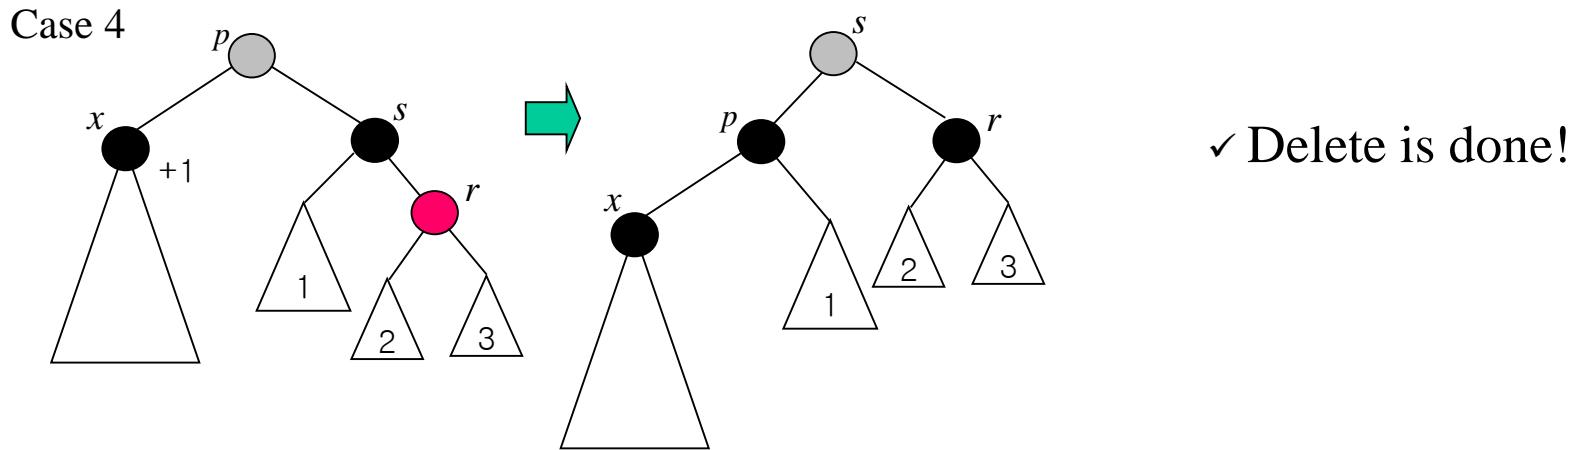


- ✓ p may have the same problem
- $O(\log n)$ time

Case 3: s is black, s 's left child is red, and s 's right child is black



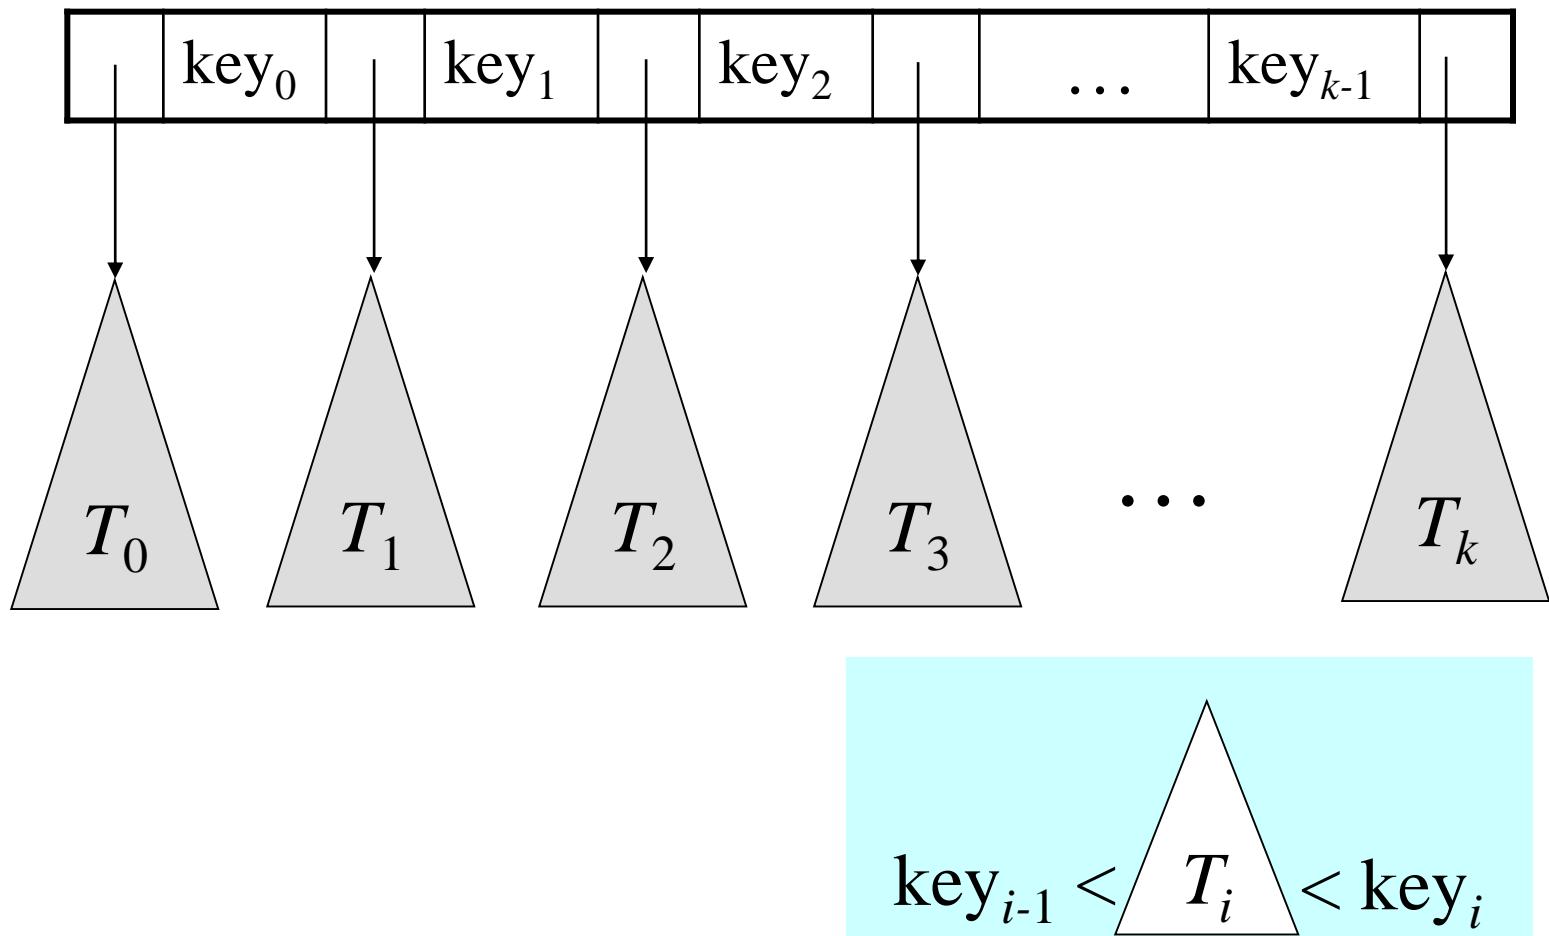
Case 4: s is black, and s 's right child is red



B-Trees

- Access unit
 - Memory to disk: page (4KB)
 - Cache to memory: cache line (64B)
- Disk access is costly (equivalent to processing time of several hundred thousand instructions).
- If a search tree is stored in a disk, minimize the height of search tree
- B-tree is a balanced multi-way search tree that minimizes disk accesses in the worst case.

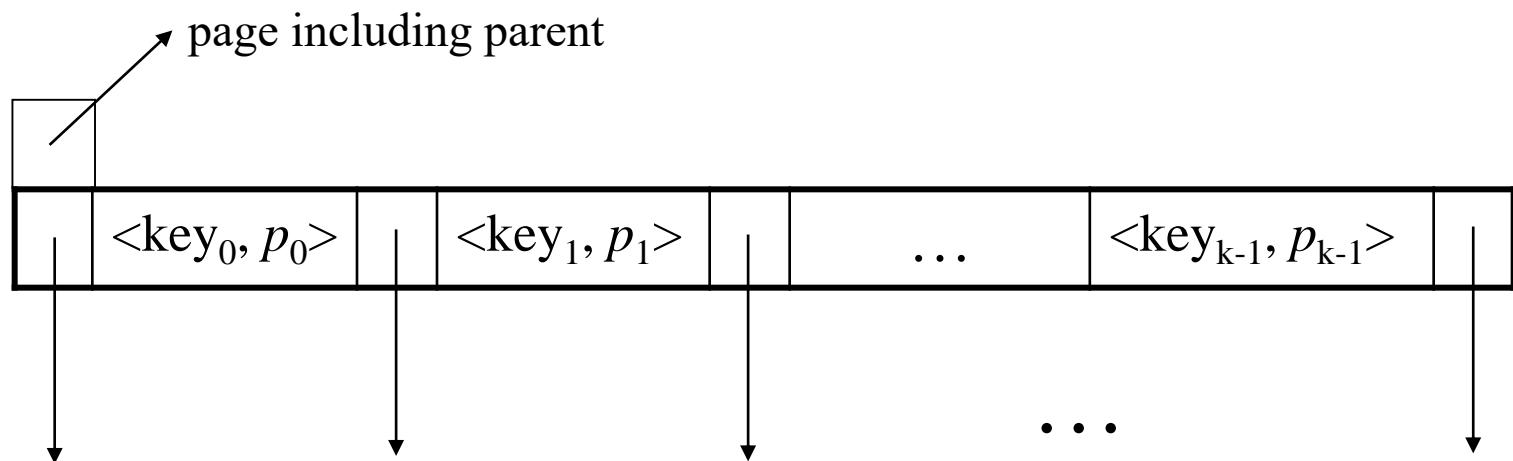
Multi-Way Search Tree



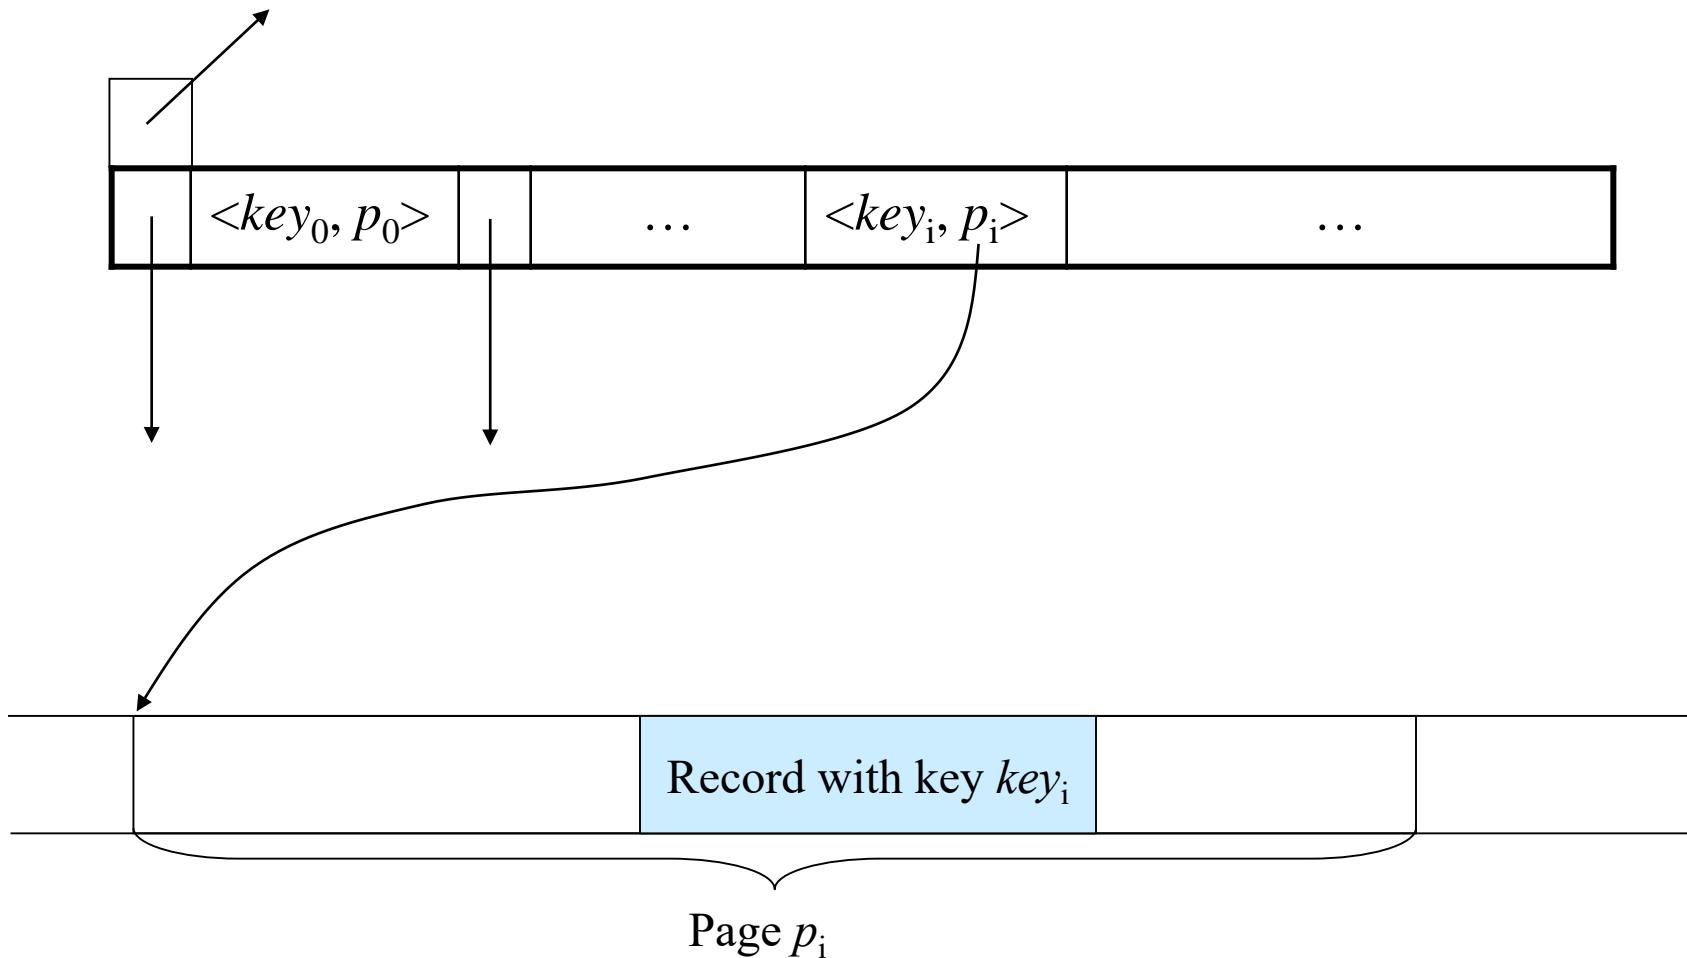
B-Tree

- B-tree is a balanced multi-way search tree that satisfies the following properties.
 - Every node except root has $\lfloor k/2 \rfloor \sim k$ keys.
 - All leaves have the same depths.

Node Structure of B-Tree



Access to Record through B-Tree



Insert in B-Tree

```
BTreeInsert( $t, x$ )  
{
```

▷ t : root node
▷ x : key to be inserted

find leaf r where x should be inserted;
insert x into r ;
if (overflow in r) **then** clearOverflow(r);

```
}
```

clearOverflow(r)
{

if (r 's sibling s has room) **then** {move key in r to s };
else {
 split r into two, and move middle key to parent p ;
 if (overflow in parent p) **then** clearOverflow(p);

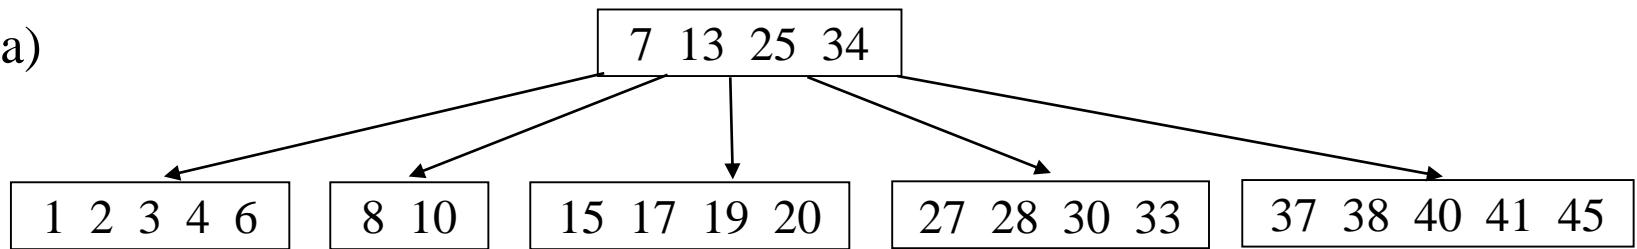
```
}
```

}

Insert in B-Tree

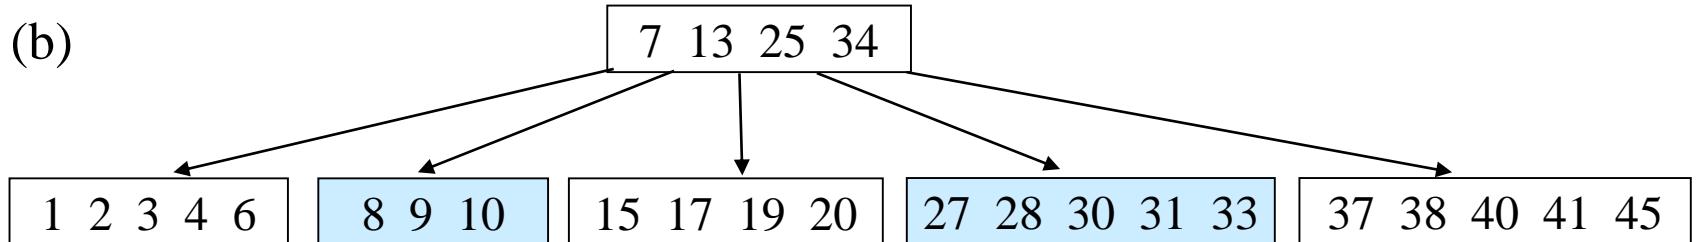
$k = 5$

(a)



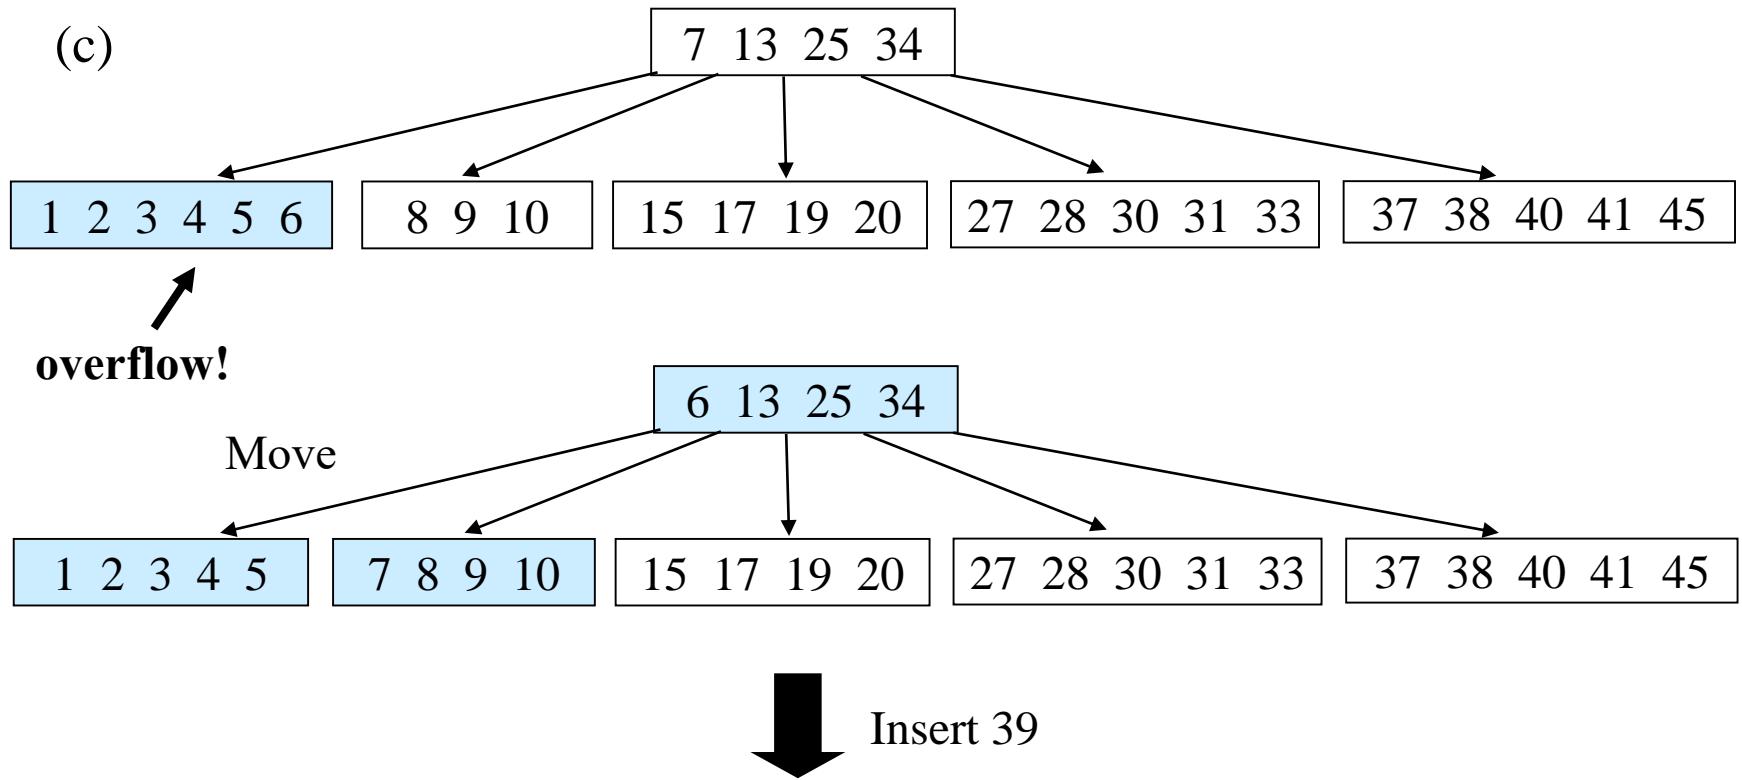
↓
Insert 9, 31

(b)



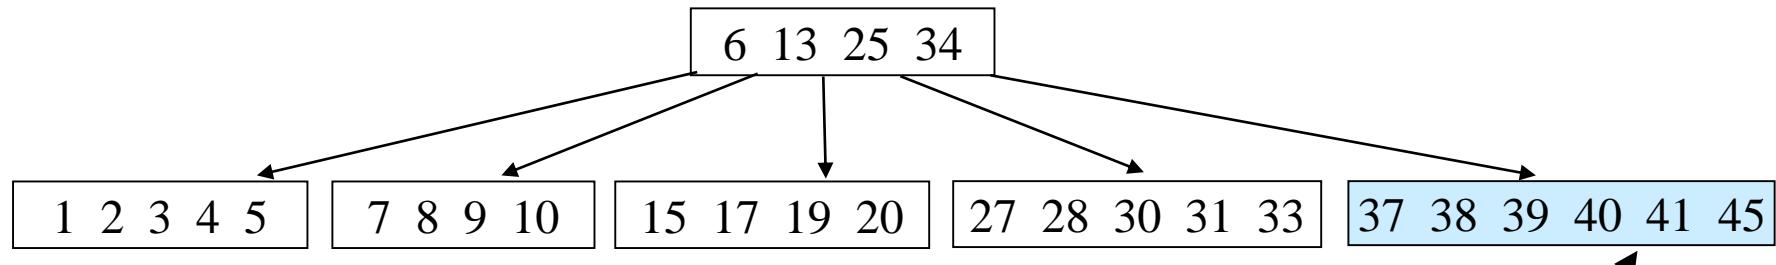
↓
Insert 5

(c)

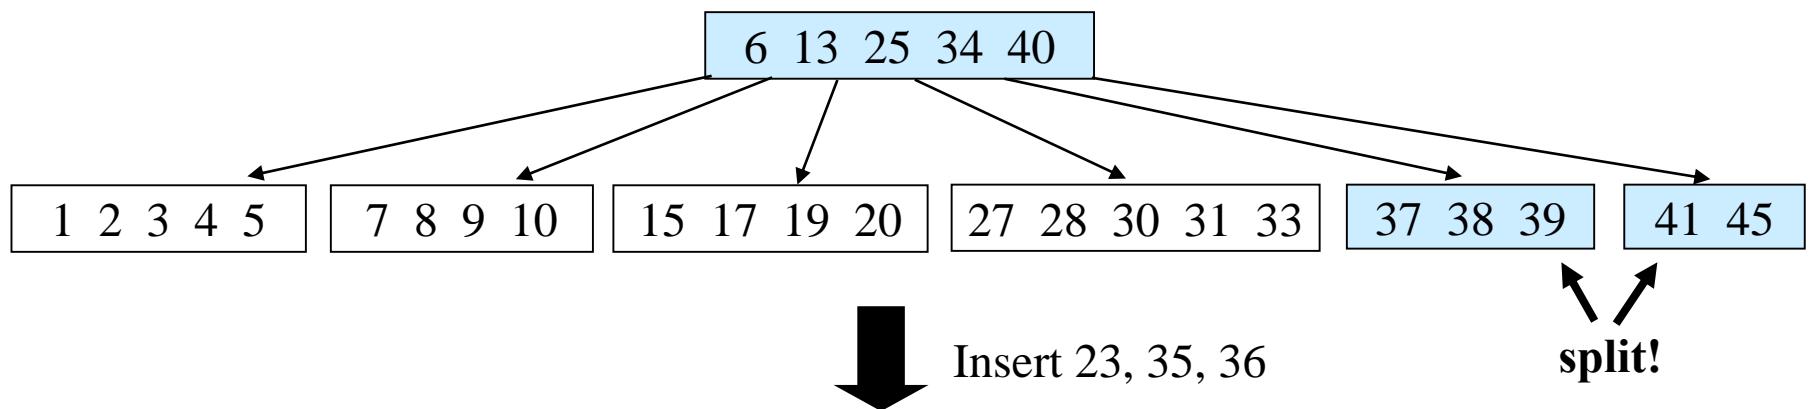


(d)

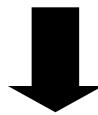
↓ Insert 39



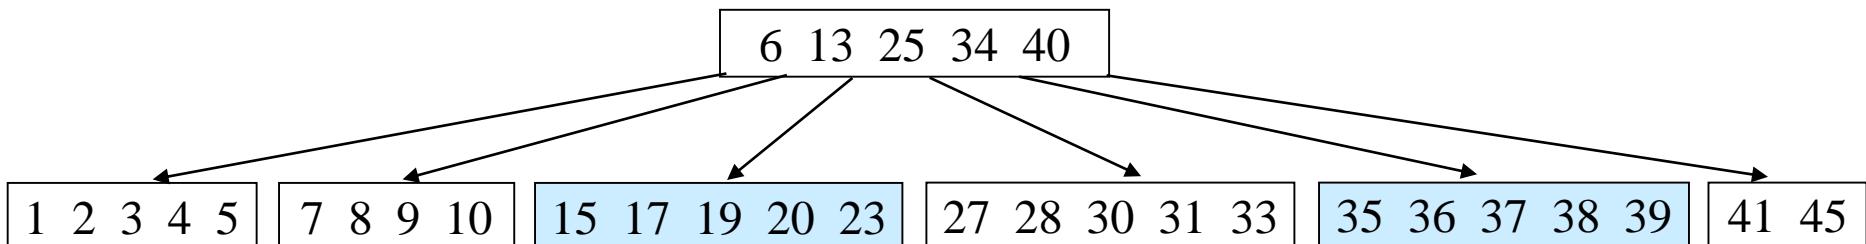
overflow!



(e)



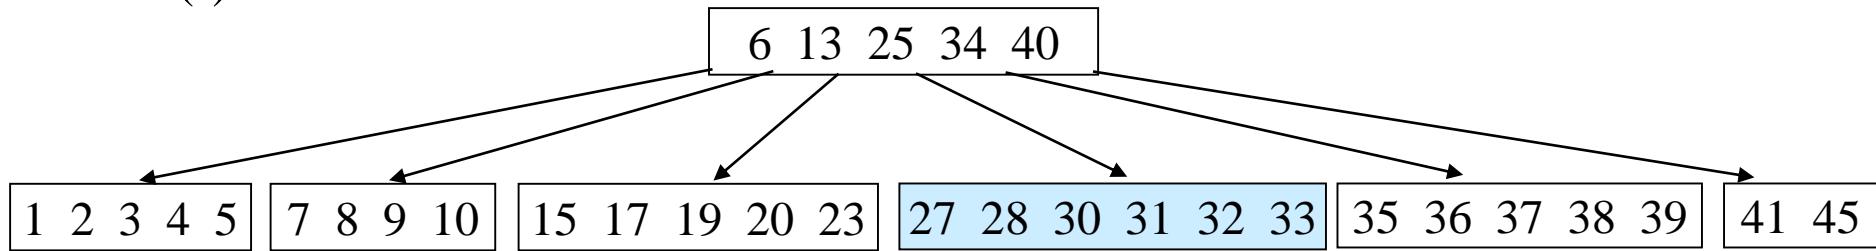
Insert 23, 35, 36



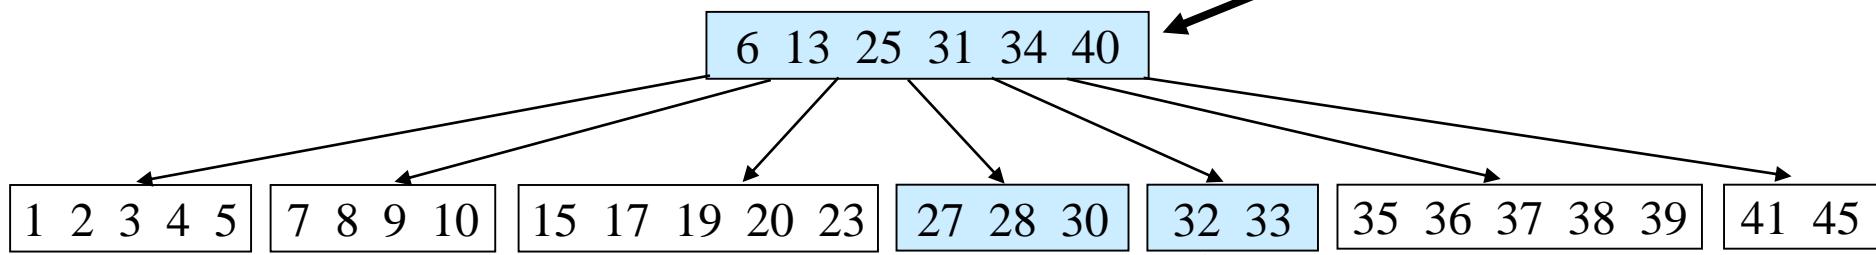
Insert 32

(f)

Insert 32

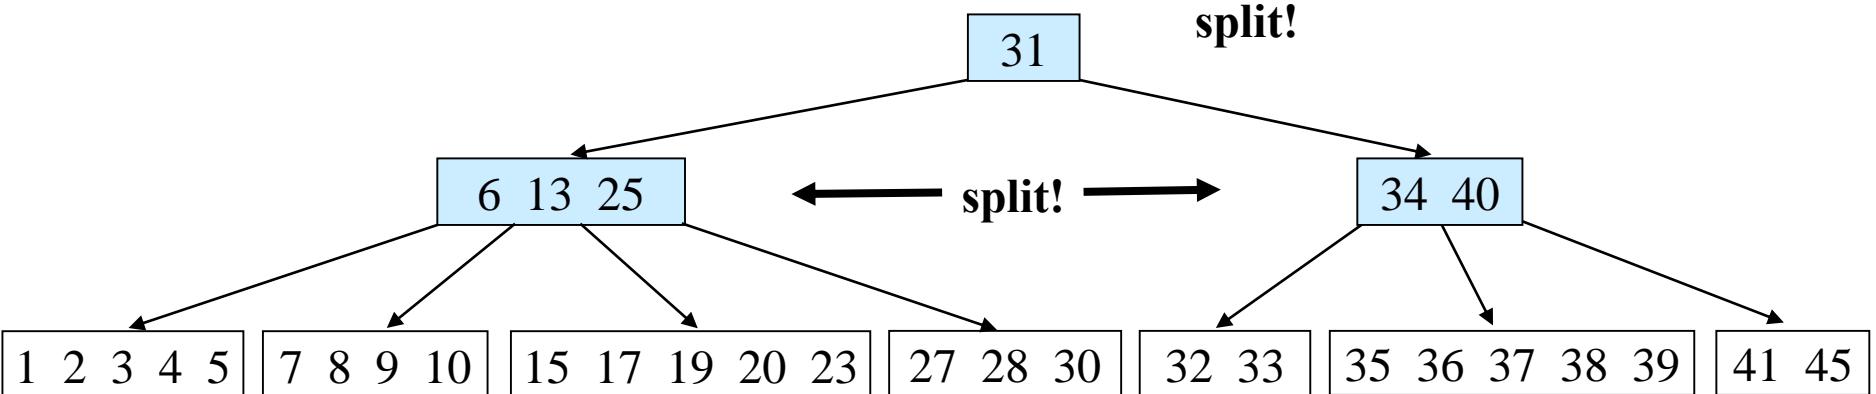


overflow!
overflow!



split!

split! ← →



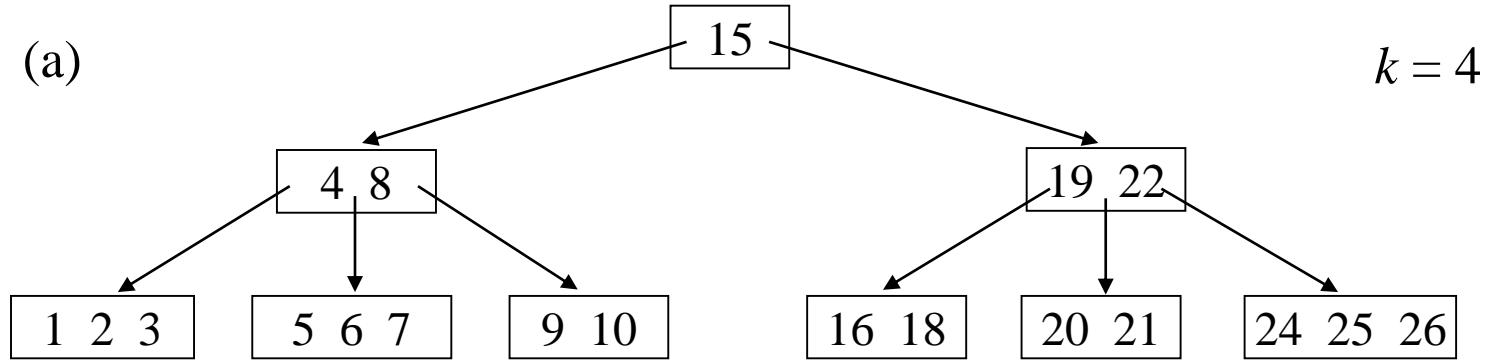
Delete in B-Tree

```
BTreeDelete( $t, x, v$ )
{
    if ( $v$  not leaf) then {
        find leaf  $r$  that contains  $x$ 's successor  $y$ ;
        swap  $x$  and  $y$ ;
    }
    delete  $x$  in  $r$ ;
    if (underflow in  $r$ ) then clearUnderflow( $r$ );
}
clearUnderflow( $r$ )
{
    if ( $r$ 's sibling  $s$  has keys to move)
        then {move key in  $s$  to  $r$ ;}
        else {
            merge  $r$  and  $s$ ;
            if (underflow in parent  $p$ ) then clearUnderflow( $p$ );
        }
}
```

- ▷ t : root node
- ▷ x : key to be deleted
- ▷ v : node containing x

Delete in B-Tree

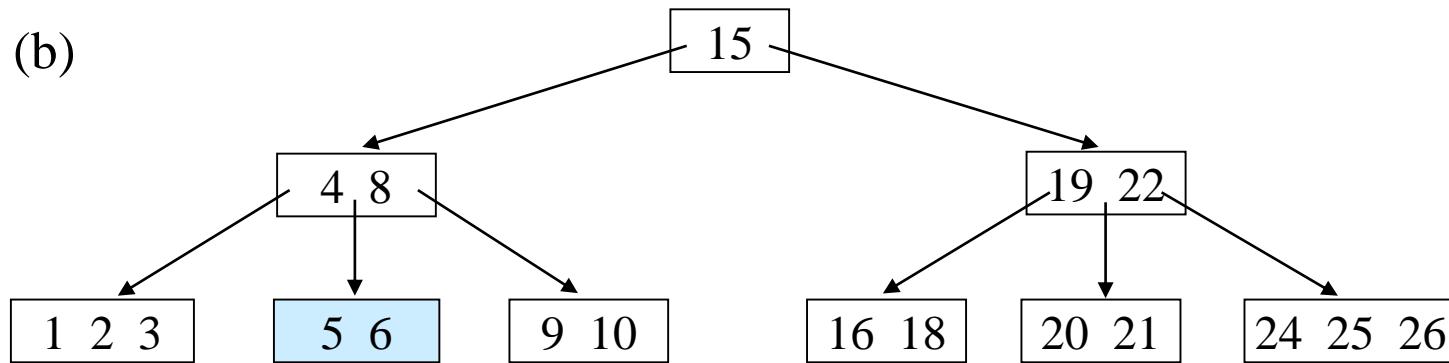
(a)



$k = 4$

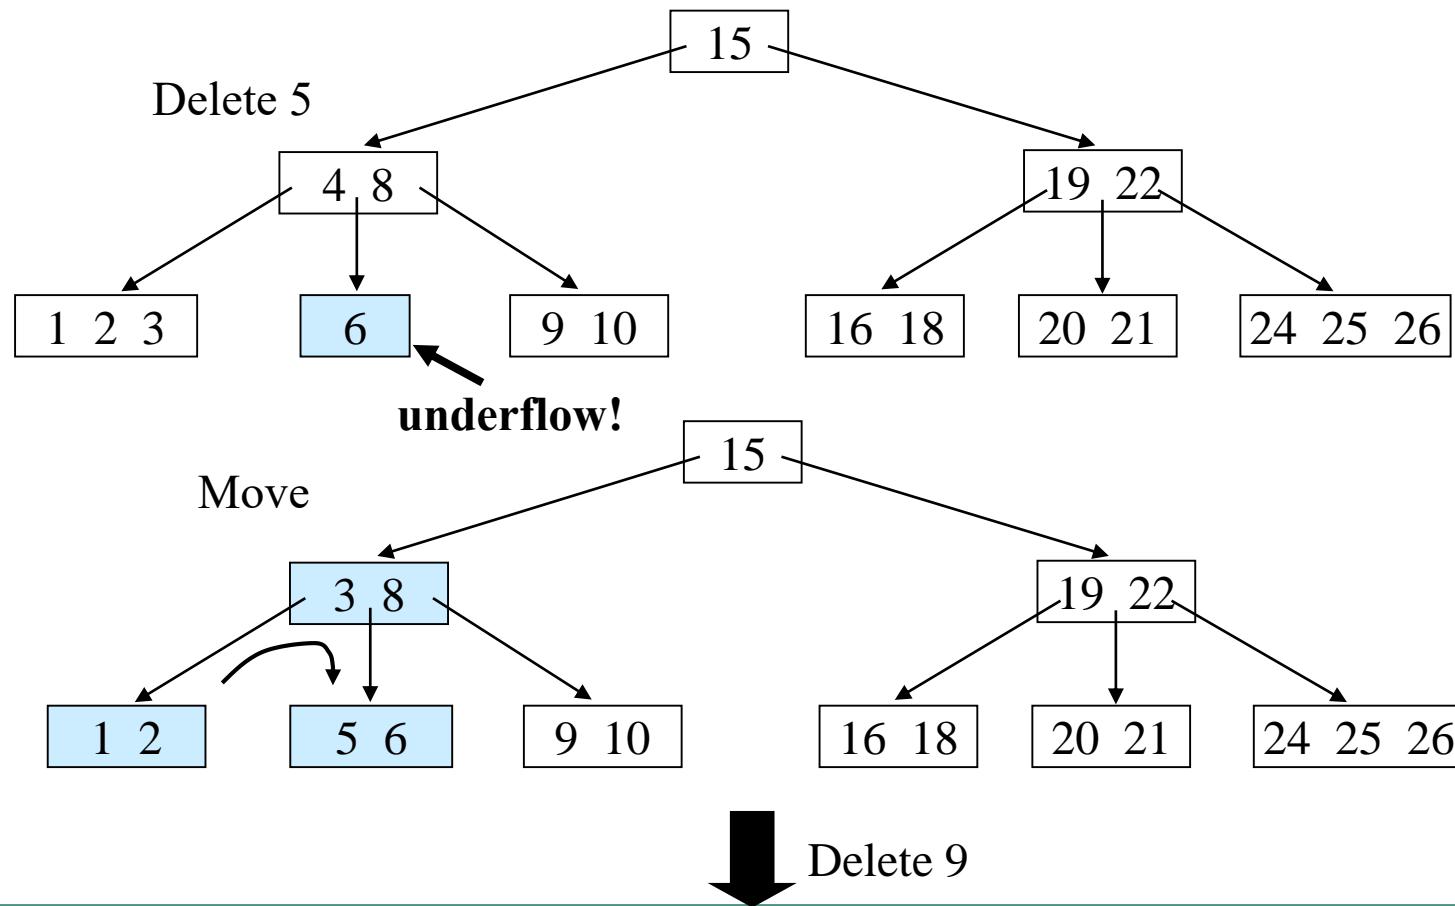
Delete 7

(b)

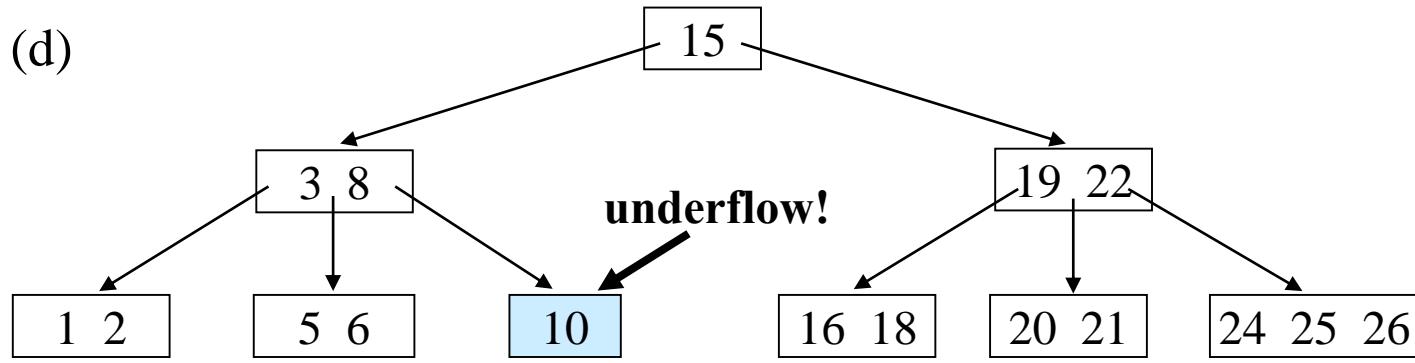


Delete 5

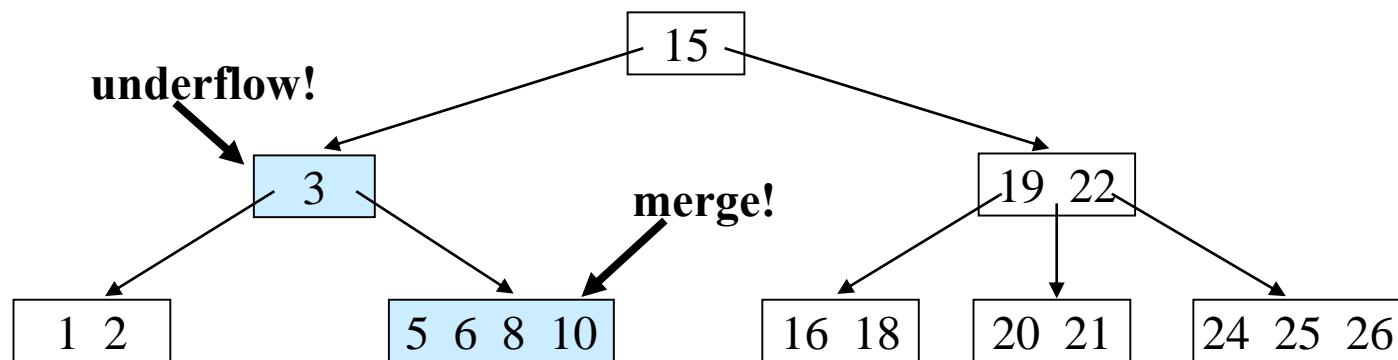
(c)



(d)



underflow!



merge!

merge!

